



# CSE 412 / 598 : Database Management

Lecture 12

Instructor: Mohamed Sarwat



# Transaction Management Overview

# Transactions

Concurrent execution of user programs is essential for good DBMS performance.

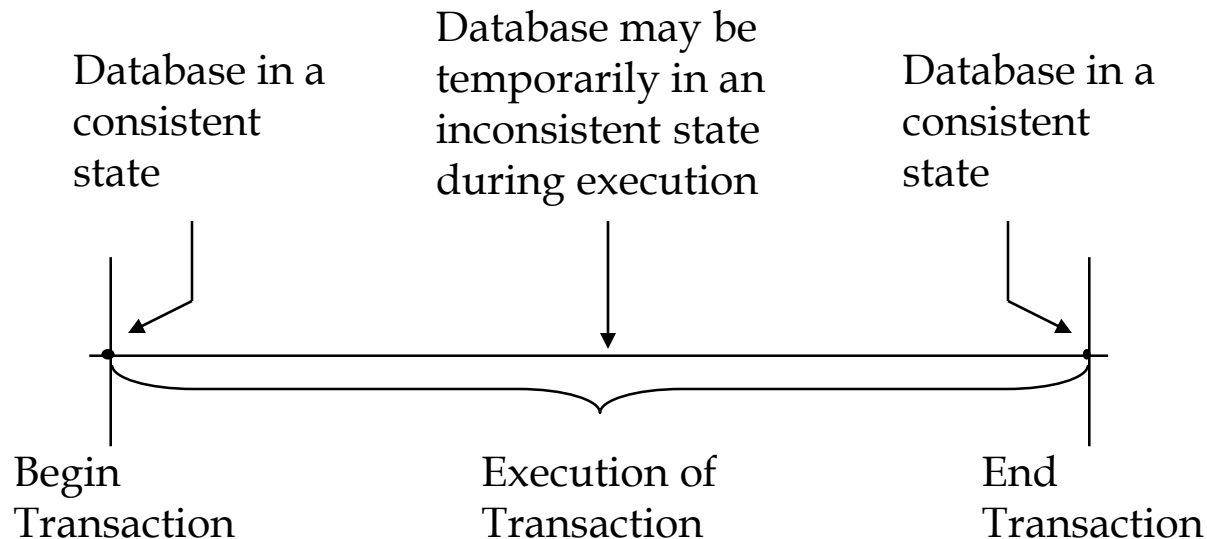
Why ?

# Transactions

- Concurrent execution of user programs is essential for good DBMS performance.
- A user's program may carry out many operations on the data retrieved from the database, but the DBMS is only concerned about what data is read/written from/to the database.
- A transaction is the DBMS's abstract view of a user program: a sequence of reads and writes.

# Transaction

A transaction is a collection of actions that make consistent transformations of system states while preserving system consistency.



# Concurrency in a DBMS

- Users submit transactions, and can think of each transaction as executing by itself.
  - Concurrency is achieved by the DBMS, which interleaves actions (reads/writes of DB objects) of various transactions.
  - Each transaction must leave the database in a consistent state if the DB is consistent when the transaction begins.
    - DBMS will enforce some ICs, depending on the ICs declared in CREATE TABLE statements.
    - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).
- Issues: Effect of *interleaving* transactions, and *crashes*.

# Principles of Transactions

## **A** TOMICITY

- all or nothing

## **C**ONSISTENCY

- no violation of integrity constraints

## **I**SOLATION

- concurrent changes invisible  $\Rightarrow$  serializable

## **D**URABILITY

- committed updates persist

# Atomicity of Transactions

- A transaction might *commit* after completing all its actions, or it could *abort* (or be aborted by the DBMS) after executing some actions.
- A very important property guaranteed by the DBMS for all transactions is that they are *atomic*. That is, a user can think of a Xact as always executing all its actions in one step, or not executing any actions at all.
  - DBMS *logs* all actions so that it can *undo* the actions of aborted transactions.



# Consistency

- Internal consistency
  - A transaction which executes **alone** against a **consistent** database leaves it in a consistent state.
  - Transactions do not violate database integrity constraints.
- Transactions are **correct** programs

# Isolation

- Serializability
  - If several transactions are executed concurrently, the results must be the same as if they were executed serially in some order.
- Incomplete results
  - An incomplete transaction cannot reveal its results to other transactions before its commitment.
  - Necessary to avoid cascading aborts.

# Durability

- Once a transaction commits, the system must guarantee that the results of its operations will never be lost, in spite of subsequent failures.
- Database recovery

# Example

- Consider two transactions (*Xacts*):

T1:	BEGIN	$A=A+100$ ,	$B=B-100$	END
T2:	BEGIN	$A=1.06*A$ ,	$B=1.06*B$	END

- ❖ Intuitively, the first transaction is transferring \$100 from B's account to A's account. The second is crediting both accounts with a 6% interest payment.
- ❖ There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together. However, the net effect *must* be equivalent to these two transactions running serially in some order.

## Example (Contd.)

- Consider a possible interleaving (*schedule*):

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A,$	$B = 1.06 * B$

- ❖ This is OK. But what about:

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A, B = 1.06 * B$	

- ❖ The DBMS' s view of the second schedule:

T1:	$R(A), W(A),$	$R(B), W(B)$
T2:	$R(A), W(A), R(B), W(B)$	

# Scheduling Transactions

- *Serial schedule*: Schedule that does not interleave the actions of different transactions.
- *Equivalent schedules*: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- *Serializable schedule*: A schedule that is equivalent to some serial execution of the transactions.

# Interleaved Execution

## S1

T1:	R(A), W(A),	R(B), W(B), Abort
T2:	R(A), W(A), C	

## S2

T1:	R(A),	R(A), W(A), C
T2:	R(A), W(A), C	

## S3

T1:	W(A),	W(B), C
T2:	W(A), W(B), C	

**Identify Anomalies ?**

# Anomalies (Continued)

- Overwriting Uncommitted Data (WW Conflicts):

T1:	W(A),	W(B), C
T2:	W(A), W(B), C	



# Anomalies with Interleaved Execution

- Reading Uncommitted Data (WR Conflicts, “dirty reads”):

T1:	R(A), W(A),	R(B), W(B), Abort
T2:	R(A), W(A), C	

- Unrepeatable Reads (RW Conflicts):

T1:	R(A),	R(A), W(A), C
T2:	R(A), W(A), Abort	

# Anomalies (Continued)

- Overwriting Uncommitted Data (WW Conflicts):

T1:	W(A),	W(B), Abort
T2:	W(A), W(B), C	

# Conflict Serializable Schedules

- Two schedules are **conflict equivalent** if:
  - Involve the same actions of the same transactions
  - Every pair of conflicting actions is ordered the same way
- Schedule S is **conflict serializable** if S is conflict equivalent to some serial schedule



T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	

Conflict Serializable?



T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A),	R(B), W(B)

**Conflict Serializable?**



T1: R(A), W(A), R(B), W(B)


T2: R(A), W(A), R(B), W(B)

**Conflict Serializable?**

# Example

- A schedule that is not conflict serializable:

T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	



- The cycle in the graph reveals the problem. The output of T1 depends on T2, and vice-versa.

# Dependency Graph

- Dependency graph: One node per  $X_{act}$ ; edge from  $T_i$  to  $T_j$  if  $T_j$  reads/writes an object last written by  $T_i$ .
- Theorem: Schedule is conflict serializable if and only if its dependency graph is acyclic



# Lock-Based Concurrency Control

- *Strict Two-phase Locking (Strict 2PL) Protocol:*
  - Each Xact must obtain a *S (shared) lock* on object before reading, and an *X (exclusive) lock* on object before writing.
  - All locks held by a transaction are released when the transaction completes
    - If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object.
- Strict 2PL allows only serializable schedules.
  - Additionally, it simplifies transaction aborts
  - *(Non-strict) 2PL* also allows only serializable schedules, but involves more complex abort processing

# Lock Management

- Lock and unlock requests are handled by the lock manager
- Lock table entry:
  - Number of transactions currently holding a lock
  - Type of lock held (shared or exclusive)
  - Pointer to queue of lock requests
- Locking and unlocking have to be atomic operations
- Lock upgrade: transaction that holds a shared lock can be upgraded to hold an exclusive lock

# Deadlocks

- Deadlock: Cycle of transactions waiting for locks to be released by each other.



# Deadlock Detection

- Create a **waits-for graph**:
  - Nodes are transactions
  - There is an edge from  $T_i$  to  $T_j$  if  $T_i$  is waiting for  $T_j$  to release a lock
- Periodically check for cycles in the waits-for graph

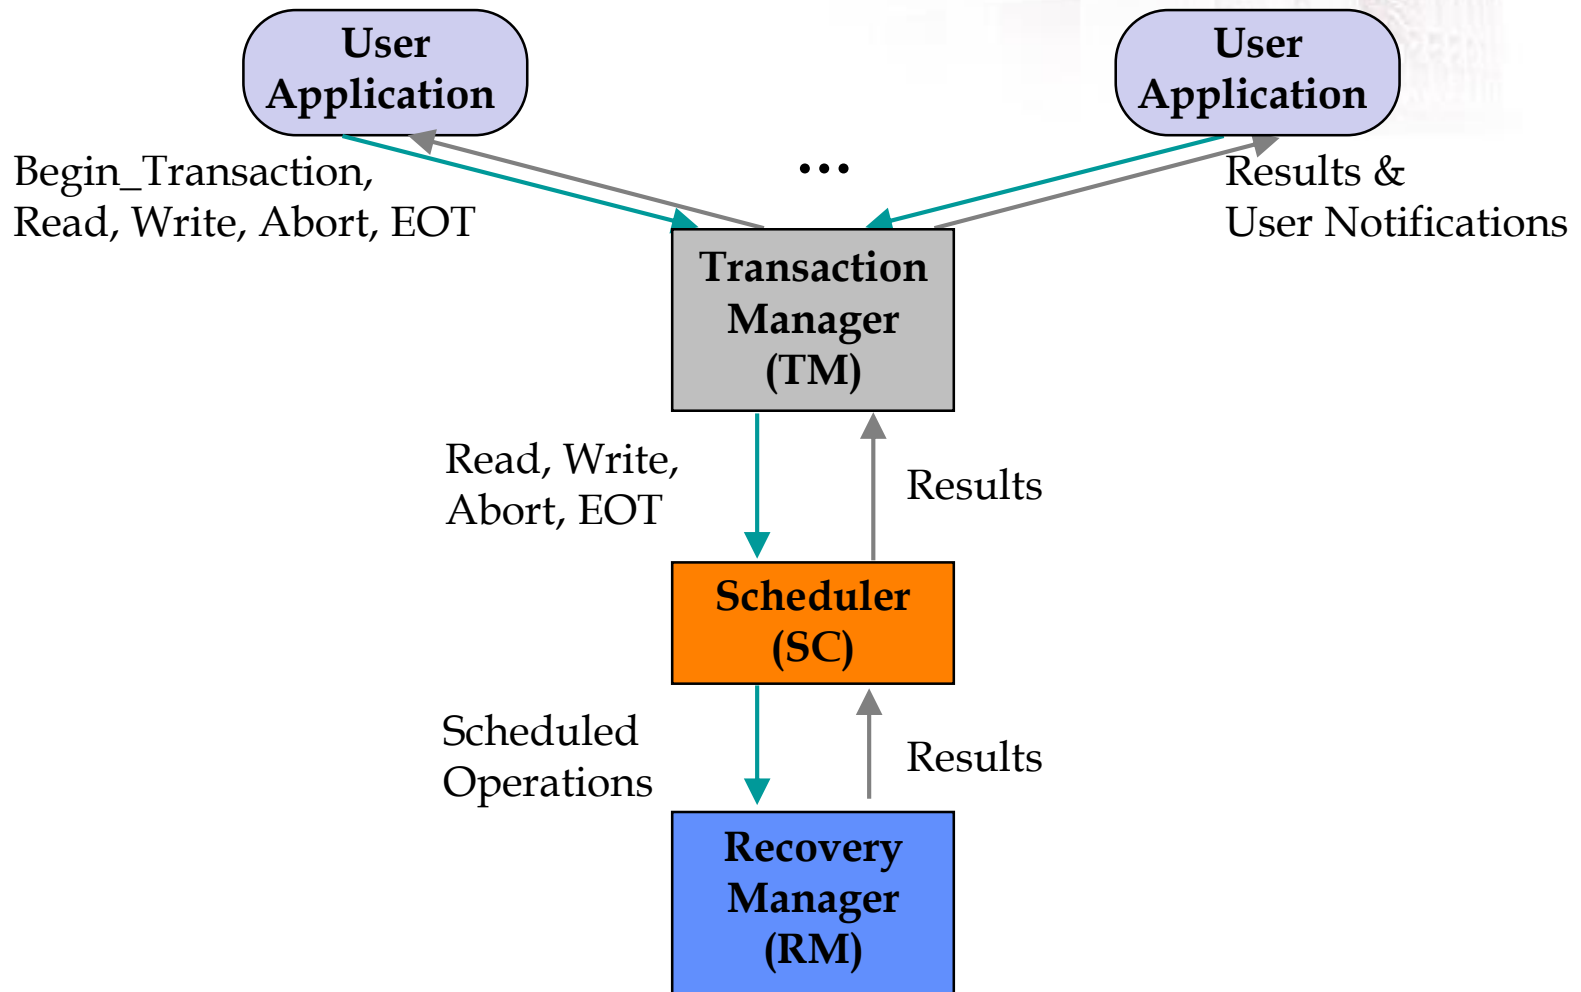
# Deadlock Detection (Continued)

Example:

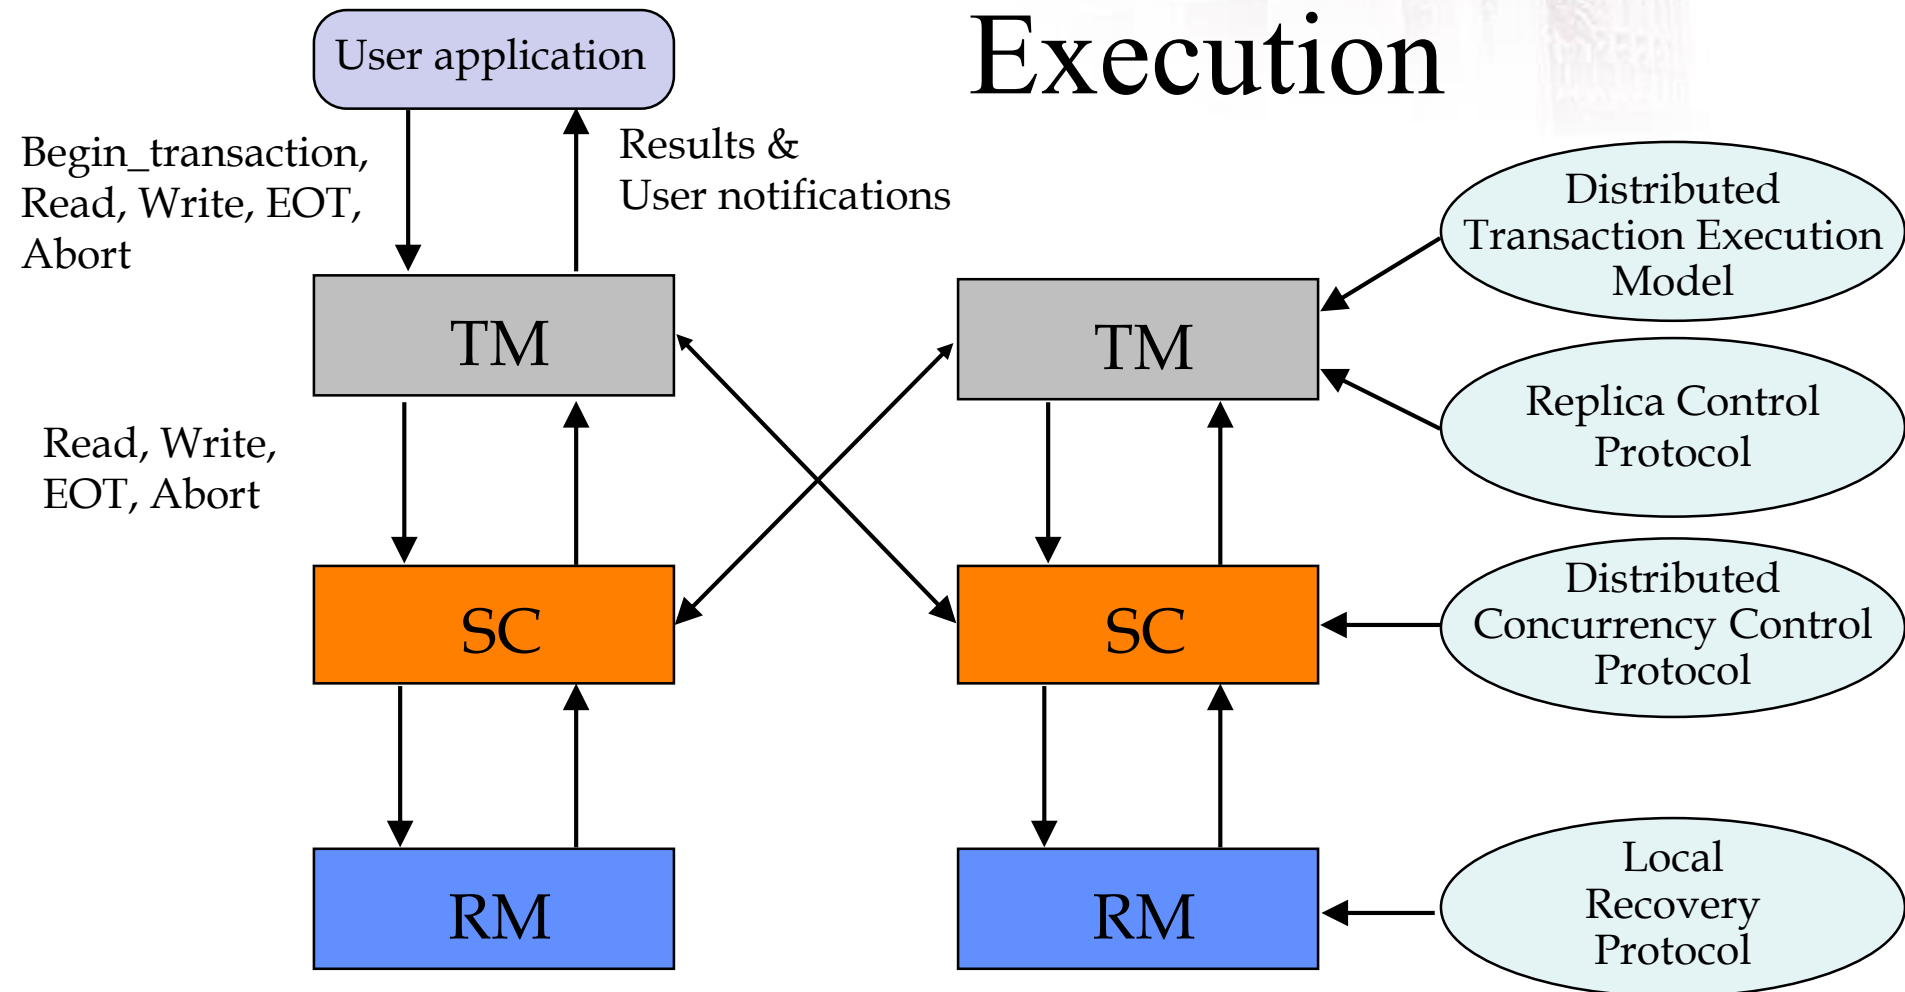
T1: S(A), R(A), S(B)  
 T2: X(B), W(B) X(C)  
 T3: S(C), R(C) X(A)  
 T4: X(B)



# Centralized Transaction Execution



# Distributed Transaction Execution





# Questions